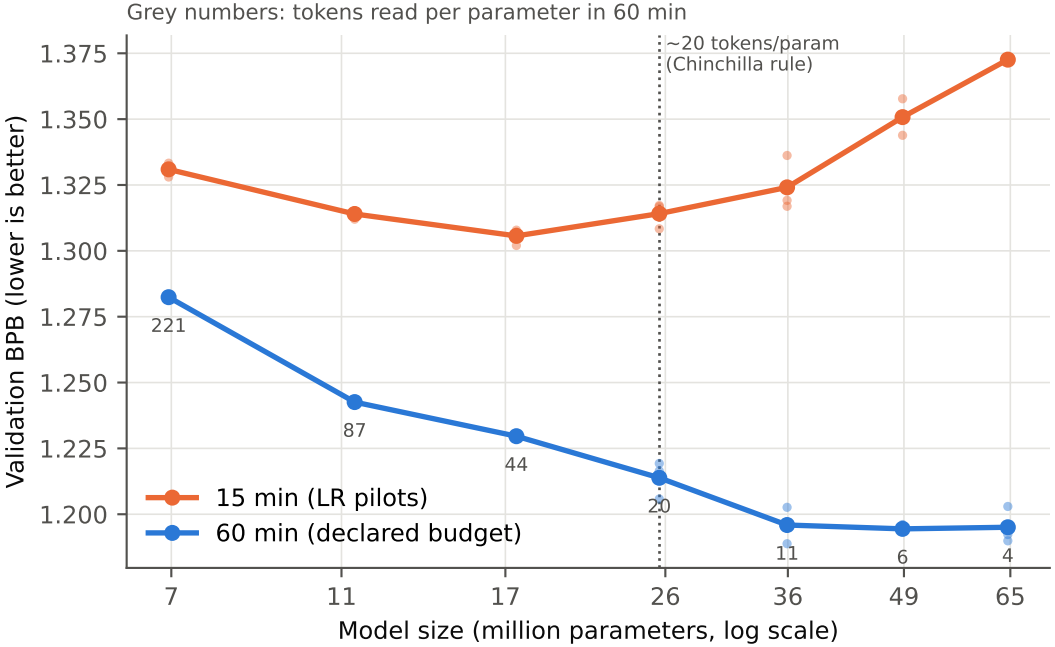
Grey numbers: tokens read per parameter in 60 min
~20 tokens/param
(Chinchilla rule)
Validation BPB (lower is better)
1.375
1.350
1.325
1.300
1.275
1.250
1.225
1.200
221
87
44
20
11
6
4
15 min (LR pilots)
60 min (declared budget)
7
11
17
26
36
49
65
Model size (million parameters, log scale)